

Final Take Home Test
Optimization of Matrix-Matrix Multiplication Using
Vector Instructions

Zi Xuan Li
Professor Izidor Gertner
CSC 34200 Spring 2024
May 12th, 2024

Table of Contents

1. Objective.....	3
2. Task#1.....	3
3. Task#2.....	3-5
4. Task#3.....	5-7
5. Task#4.....	8
6. Task#5.....	8-12
7. Task#6.....	12-14
8. Task#7.....	14-16
9. Conclusion.....	16

Objective:

The aim of this concluding take-home exam is to enhance compiler-generated code for computing the product of two matrices using vector instructions and DPPS vector instruction. It builds upon the earlier test on matrix multiplication. The objective is to explore the performance contrast between vectorization and non-vectorization, assessed with Chrono, a high-resolution timer for execution time measurement.

Task#1: Use CPUID instructions to determine your processor vector processing capabilities.

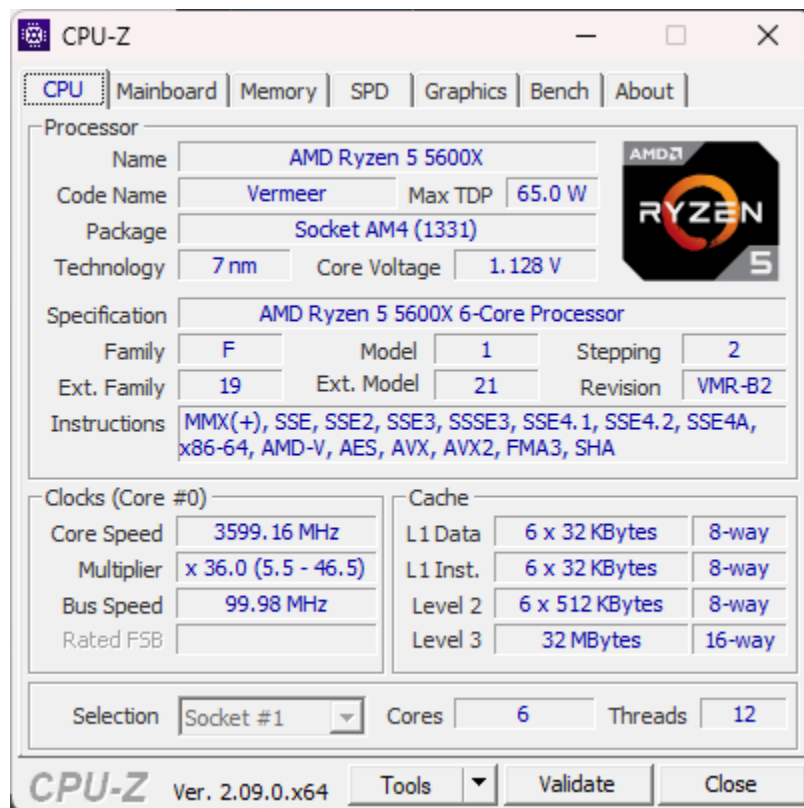
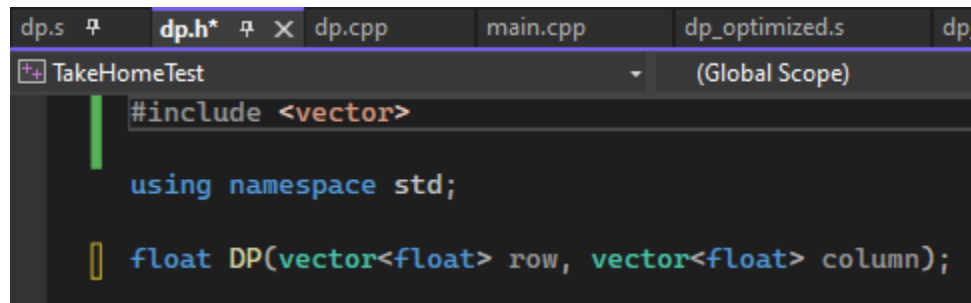


Figure 1: Information of my CPU using CPU-Z from the CPUID instructions

The processing capabilities of my CPU, the AMD Ryzen 5 5600x, can be found under the instructions tab in the window above. It supports MMX (+), SSE, SSE2, SSE3, SSSE3, SSE4.1, SSE4.2, SSE4A, x86-64, AMD-V, AES, AVX, AVX2, FMA3, and SHA.

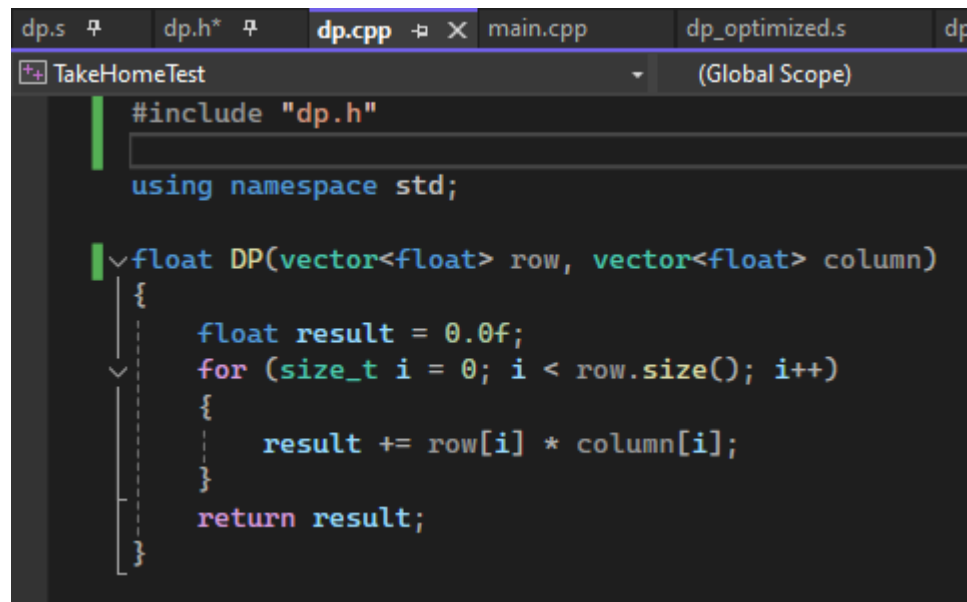
Task#2: Write C/C++ main () to compute Matrix-Matrix multiplication, can be taken from previous take-home test. The focus in this take home test is the most inner loop that computes dot product of row and column. You can use a function DP (row, column) from a previous take home test.

Place the function DP (row, column) in a separate file from main () that calls this function. Vector sizes should be powers of 2 (e.g. 16, 32, 64,512, ...216 etc.)



```
dp.s  dp.h*  dp.cpp  main.cpp  dp_optimized.s  dp_  
TakeHomeTest  (Global Scope)  
#include <vector>  
using namespace std;  
float DP(vector<float> row, vector<float> column);
```

Figure 2: Image of the dp.h file



```
dp.s  dp.h*  dp.cpp  main.cpp  dp_optimized.s  dp_  
TakeHomeTest  (Global Scope)  
#include "dp.h"  
using namespace std;  
float DP(vector<float> row, vector<float> column)  
{  
    float result = 0.0f;  
    for (size_t i = 0; i < row.size(); i++)  
    {  
        result += row[i] * column[i];  
    }  
    return result;  
}
```

Figure 3: Image of the dp.cpp file

```

dp.s  dp.h  dp.cpp  main.cpp  dp_optimized.s  dp_dpps.cpp
TakeHomeTest (Global Scope)

#include "dp.h"
#include <iostream>
#include <vector>
#include <stdio.h>
#include <random>
#include <chrono>

using namespace std;

void fill_matrix(vector<float>& matrix, size_t N) {
    random_device rd;
    mt19937 gen(rd());
    uniform_real_distribution<> dist(0.0f, 99.0f);
    for (size_t i = 0; i < N * N; ++i) {
        matrix[i] = dist(gen);
    }
}

void matrix_multiply(const vector<float>& A, const vector<float>& B, vector<float>& C, size_t N) {
    for (size_t i = 0; i < N; ++i) {
        for (size_t j = 0; j < N; ++j) {
            vector<float> row(N), column(N);
            for (size_t k = 0; k < N; ++k) {
                row[k] = A[i * N + k];
                column[k] = B[k * N + j];
            }
            C[i * N + j] = DP(row, column);
        }
    }
}

int main() {
    vector<size_t> sizes = { 16, 32, 64, 128, 256, 512, 1024 };
    for (auto N : sizes) {
        vector<float> A(N * N), B(N * N), C(N * N);

        fill_matrix(A, N);
        fill_matrix(B, N);

        auto start = chrono::high_resolution_clock::now();
        matrix_multiply(A, B, C, N);
        auto end = chrono::high_resolution_clock::now();
        chrono::duration<double> diff = end - start;
        cout << "Time to multiply two " << N << "x" << N << " float matrices: " << diff.count() << " seconds\n";
    }
    return 0;
}

```

Figure 4: Image of the main.cpp file.

Task#3: Compile code in §2 and create assembly code for function DP (row, column) only. Make sure that compiler generated vectorized code.

```

dp.h  dp.cpp  main.cpp*  dp.s  dp_optimized.s  dp_dpps.cpp
1      .file "dp.cpp"
2      .text
3      .p2align 4
4      .globl _Z2DPSt6vectorIfSaIfEES1_
5      .def _Z2DPSt6vectorIfSaIfEES1_; .scl 2; .type 32; .endef
6      .seh_proc _Z2DPSt6vectorIfSaIfEES1_
7      _Z2DPSt6vectorIfSaIfEES1_:
8      .LFB1021:
9      .seh_endprologue
10     movq 8(%rcx), %r8
11     movq (%rcx), %rax
12     movq %r8, %r9
13     subq %rax, %r9
14     movq %r9, %rcx
15     sarq $2, %rcx
16     cmpq %rax, %r8
17     je .L10
18     movq (%rdx), %rdx
19     cmpq $28, %r9
20     jbe .L11
21     movq %rcx, %r9
22     xorl %r8d, %r8d
23     vxorps %xmm0, %xmm0, %xmm0
24     shrq $3, %r9
25     salq $5, %r9
26     .p2align 4
27     .p2align 3
28     .L4:
29     vmovups (%rax,%r8), %ymm4
30     vmulps (%rdx,%r8), %ymm4, %ymm1
31     addq $32, %r8
32     vshufps $85, %xmm1, %xmm1, %xmm3
33     vshufps $255, %xmm1, %xmm1, %xmm2
34     vaddss %xmm1, %xmm0, %xmm0
35     vaddss %xmm3, %xmm0, %xmm0
36     vunpckhps %xmm1, %xmm1, %xmm3
37     vextractf128 $0x1, %ymm1, %xmm1
38     vaddss %xmm3, %xmm0, %xmm0
39     vaddss %xmm2, %xmm0, %xmm0
40     vshufps $85, %xmm1, %xmm1, %xmm2
41     vaddss %xmm1, %xmm0, %xmm0
42     vaddss %xmm2, %xmm0, %xmm0
43     vunpckhps %xmm1, %xmm1, %xmm2
44     vshufps $255, %xmm1, %xmm1, %xmm1
45     vaddss %xmm2, %xmm0, %xmm0
46     vaddss %xmm1, %xmm0, %xmm0
47     cmpq %r8, %r9
48     jne .L4
49     movq %rcx, %r8
50     andq $-8, %r8
51     testb $7, %cl

```

```

dp.h  dp.cpp  main.cpp*  dp.s  dp_optimized.s  dp_dpps.cpp
51      testb  $7, %cl
52      je    .L22
53      vzeroupper
54  ~.L3:
55      movq   %rcx, %r9
56      subq   %r8, %r9
57      leaq   -1(%r9), %r10
58      cmpq   $2, %r10
59      jbe    .L9
60      vmovups (%rax,%r8,4), %xmm5
61      movq   %r9, %r10
62      vmulps (%rdx,%r8,4), %xmm5, %xmm1
63      andq   $-4, %r10
64      addq   %r10, %r8
65      andl   $3, %r9d
66      vaddss %xmm1, %xmm0, %xmm0
67      vshufps $85, %xmm1, %xmm1, %xmm2
68      vaddss %xmm2, %xmm0, %xmm0
69      vunpckhps %xmm1, %xmm1, %xmm2
70      vshufps $255, %xmm1, %xmm1, %xmm1
71      vaddss %xmm2, %xmm0, %xmm0
72      vaddss %xmm1, %xmm0, %xmm0
73      je     .L1
74      .p2align 4
75      .p2align 3
76  ~.L9:
77      vmovss (%rax,%r8,4), %xmm1
78      vmulss (%rdx,%r8,4), %xmm1, %xmm1
79      incq   %r8
80      vaddss %xmm1, %xmm0, %xmm0
81      cmpq   %rcx, %r8
82      jnb    .L9
83  ~.L1:
84      ret
85      .p2align 4
86      .p2align 3
87  ~.L10:
88      vxorps %xmm0, %xmm0, %xmm0
89      ret
90      .p2align 4
91      .p2align 3
92  ~.L22:
93      vzeroupper
94      ret
95  ~.L11:
96      xorl   %r8d, %r8d
97      vxorps %xmm0, %xmm0, %xmm0
98      jmp    .L3
99      .seh_endproc
100     .ident  "GCC: (Rev2, Built by MSYS2 project) 12.1.0"
101

```

Figure 5: Image of the compiler generated assembly code from dp.h

Task#4: Use high resolution timer to measure execution time (as in previous take-home test). Plot graph: time versus vector size.

```
Time to multiply two 16x16 float matrices: 9.2e-05 seconds
Time to multiply two 32x32 float matrices: 0.0002864 seconds
Time to multiply two 64x64 float matrices: 0.0009116 seconds
Time to multiply two 128x128 float matrices: 0.0057554 seconds
Time to multiply two 256x256 float matrices: 0.0376169 seconds
Time to multiply two 512x512 float matrices: 0.395068 seconds
Time to multiply two 1024x1024 float matrices: 4.69744 seconds
```

Figure 6: Execution times of various 2^n integer matrix sizes from $n = 16$ to $n = 1024$




Figure 7: Graph measuring execution time vs. vector size for compiler generated code

Task#5: Create assembly code for function DP (row, column). The assembly code should contain vector instructions.


```

dp.h  dp.cpp  main.cpp  dp.s  dp_optimized.s  dp_dpps.cpp
1      .file    "dp.cpp"
2      .text
3      .p2align 4
4      .globl  _Z2DPSt6vectorIfSaIfEES1_
5      .def    _Z2DPSt6vectorIfSaIfEES1_; .scl    2; .type  32; .endif
6      .seh_proc _Z2DPSt6vectorIfSaIfEES1_
7      _Z2DPSt6vectorIfSaIfEES1_:
8      .LFB1021:
9      .seh_endprologue
10     movq    8(%rcx), %rax
11     movq    (%rcx), %rcx
12     movq    %rax, %r9
13     subq    %rcx, %r9
14     movq    %r9, %r8
15     sarq    $2, %r8
16     cmpq    %rcx, %rax
17     je      .L9
18     movq    (%rdx), %rdx
19     cmpq    $28, %r9
20     jbe     .L10
21     movq    %r8, %r10
22     xorl    %eax, %eax
23     vxorps  %xmm0, %xmm0, %xmm0
24     shrq    $3, %r10
25     movq    %r10, %r11
26     salq    $5, %r11
27     andl    $1, %r10d
28     je      .L4
29     vmovups (%rcx), %ymm4
30     movl    $32, %eax
31     vmulps  (%rdx), %ymm4, %ymm1
32     vshufps $85, %xmm1, %xmm1, %xmm5
33     vunpckhps %xmm1, %xmm1, %xmm3
34     vshufps $255, %xmm1, %xmm1, %xmm2
35     vaddss  %xmm1, %xmm0, %xmm0
36     vextractf128 $0x1, %ymm1, %xmm1
37     vaddss  %xmm5, %xmm0, %xmm4
38     vaddss  %xmm3, %xmm4, %xmm0
39     vaddss  %xmm2, %xmm0, %xmm5
40     vshufps $85, %xmm1, %xmm1, %xmm0
41     vunpckhps %xmm1, %xmm1, %xmm2
42     vaddss  %xmm1, %xmm5, %xmm3
43     vshufps $255, %xmm1, %xmm1, %xmm1
44     vaddss  %xmm0, %xmm3, %xmm5
45     vaddss  %xmm2, %xmm5, %xmm4
46     vaddss  %xmm1, %xmm4, %xmm0
47     cmpq    $32, %r11
48     je      .L23
49     .p2align 4
50     .p2align 3

```

```

dp.h  dp.cpp  main.cpp  dp.s  dp_optimized.s  dp_dpps.cpp
50  | .p2align 3
51  | .L4:
52  |     vmovups (%rcx,%rax), %ymm3
53  |     vmulps (%rdx,%rax), %ymm3, %ymm5
54  |     vshufps $85, %xmm5, %xmm5, %xmm1
55  |     vunpckhps %xmm5, %xmm5, %xmm3
56  |     vshufps $255, %xmm5, %xmm5, %xmm2
57  |     vaddss %xmm5, %xmm0, %xmm0
58  |     vextractf128 $0x1, %ymm5, %xmm5
59  |     vaddss %xmm1, %xmm0, %xmm4
60  |     vmovups 32(%rcx,%rax), %ymm1
61  |     vaddss %xmm3, %xmm4, %xmm0
62  |     vaddss %xmm2, %xmm0, %xmm4
63  |     vshufps $85, %xmm5, %xmm5, %xmm0
64  |     vunpckhps %xmm5, %xmm5, %xmm2
65  |     vaddss %xmm5, %xmm4, %xmm3
66  |     vshufps $255, %xmm5, %xmm5, %xmm5
67  |     vaddss %xmm0, %xmm3, %xmm4
68  |     vaddss %xmm2, %xmm4, %xmm3
69  |     vmulps 32(%rdx,%rax), %ymm1, %ymm4
70  |     addq $64, %rax
71  |     vaddss %xmm5, %xmm3, %xmm0
72  |     vaddss %xmm4, %xmm0, %xmm5
73  |     vshufps $85, %xmm4, %xmm4, %xmm0
74  |     vunpckhps %xmm4, %xmm4, %xmm3
75  |     vshufps $255, %xmm4, %xmm4, %xmm2
76  |     vextractf128 $0x1, %ymm4, %xmm4
77  |     vaddss %xmm0, %xmm5, %xmm1
78  |     vaddss %xmm3, %xmm1, %xmm5
79  |     vaddss %xmm2, %xmm5, %xmm0
80  |     vshufps $85, %xmm4, %xmm4, %xmm5
81  |     vunpckhps %xmm4, %xmm4, %xmm2
82  |     vaddss %xmm4, %xmm0, %xmm3
83  |     vshufps $255, %xmm4, %xmm4, %xmm4
84  |     vaddss %xmm5, %xmm3, %xmm0
85  |     vaddss %xmm2, %xmm0, %xmm3
86  |     vaddss %xmm4, %xmm3, %xmm0
87  |     cmpq %rax, %r11
88  |     jne .L4
89  | .L23:
90  |     movq %r8, %r11
91  |     andq $-8, %r11
92  |     testb $7, %r8b
93  |     je .L26
94  |     vzeroupper
95  | .L3:
96  |     movq %r8, %r9
97  |     subq %r11, %r9
98  |     leaq -1(%r9), %r10
99  |     cmpq $2, %r10

```

```

dp.h 7 dp.cpp main.cpp dp.s dp_optimized.s X dp_dpps.cpp
100     jbe .L7
101     vmovups (%rcx,%r11,4), %xmm1
102     movq    %r9, %rax
103     vmulps  (%rdx,%r11,4), %xmm1, %xmm5
104     andq    $-4, %rax
105     addq    %rax, %r11
106     andl    $3, %r9d
107     vaddss  %xmm5, %xmm0, %xmm0
108     vshufps $85, %xmm5, %xmm5, %xmm3
109     vunpckhps %xmm5, %xmm5, %xmm1
110     vshufps $255, %xmm5, %xmm5, %xmm5
111     vaddss  %xmm3, %xmm0, %xmm4
112     vaddss  %xmm1, %xmm4, %xmm2
113     vaddss  %xmm5, %xmm2, %xmm0
114     je     .L1
115     v.L7:
116     vmovss  (%rcx,%r11,4), %xmm3
117     leaq    1(%r11), %r10
118     leaq    0(%r11,4), %r9
119     vfmadd231ss (%rdx,%r11,4), %xmm3, %xmm0
120     cmpq    %r8, %r10
121     jnb     .L1
122     vmovss  4(%rcx,%r9), %xmm4
123     addq    $2, %r11
124     vfmadd231ss 4(%rdx,%r9), %xmm4, %xmm0
125     cmpq    %r8, %r11
126     jnb     .L1
127     vmovss  8(%rcx,%r9), %xmm1
128     vfmadd231ss 8(%rdx,%r9), %xmm1, %xmm0
129     v.L1:
130     ret
131     .p2align 4
132     .p2align 3
133     v.L9:
134     vxorps  %xmm0, %xmm0, %xmm0
135     ret
136     .p2align 4
137     .p2align 3
138     v.L26:
139     vzeroupper
140     ret
141     v.L10:
142     xorl    %r11d, %r11d
143     vxorps  %xmm0, %xmm0, %xmm0
144     jmp     .L3
145     .seh_endproc
146     .ident  "GCC: (Rev2, Built by MSYS2 project) 12.1.0"
147

```

Figure 8: Image of the optimized assembly code of dp.h

```

Time to multiply two 16x16 float matrices: 9.12e-05 seconds
Time to multiply two 32x32 float matrices: 0.0002748 seconds
Time to multiply two 64x64 float matrices: 0.0009819 seconds
Time to multiply two 128x128 float matrices: 0.0056039 seconds
Time to multiply two 256x256 float matrices: 0.0373567 seconds
Time to multiply two 512x512 float matrices: 0.392527 seconds
Time to multiply two 1024x1024 float matrices: 4.4821 seconds

```

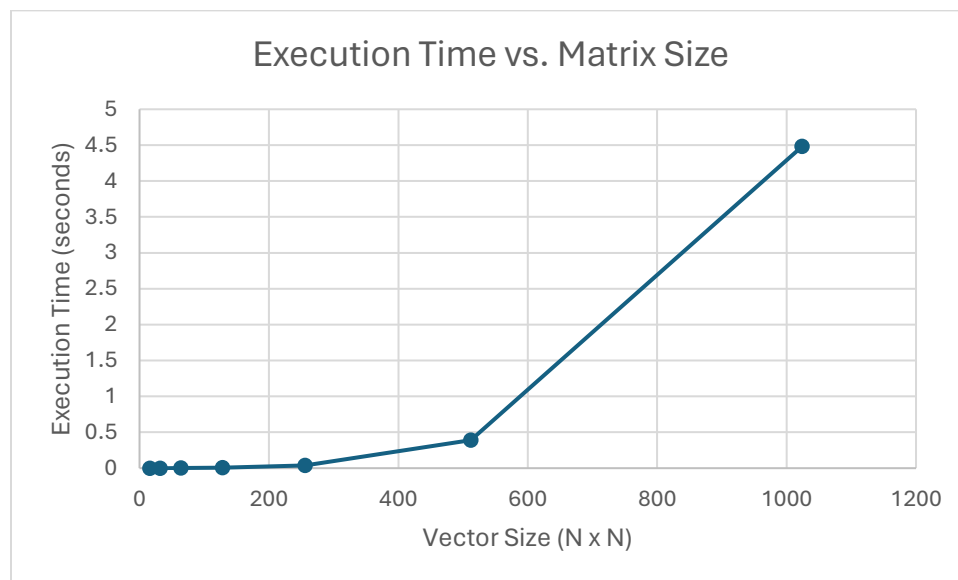
Figure 9: Execution times of various 2^n integer matrix sizes from $n = 16$ to $n = 1024$ 

Figure 10: Graph measuring execution time vs. vector size for optimized code

Task#6: To optimize the code further, please try to use machine vector instruction DPPS to compute dot product.

```

dp.h  dp.cpp  main.cpp  dp.s  dp_optimized.s  dp_dpps.cpp*
TakeHomeTest
#include "dp.h"
#include <immintrin.h>
#include <stdexcept>

using namespace std;

float DP(vector<float> row, vector<float> column)
{
    size_t size = row.size();
    if (size != column.size())
        throw std::invalid_argument("Vectors are not of the same size.");

    float result = 0.0f;
    size_t i = 0;

    __m256 sum_vec = _mm256_setzero_ps();
    for (; i <= size - 8; i += 8)
    {
        __m256 row_vec = _mm256_loadu_ps(&row[i]);
        __m256 col_vec = _mm256_loadu_ps(&column[i]);
        __m256 dp = _mm256_dp_ps(row_vec, col_vec, 0xF1);
        sum_vec = _mm256_add_ps(sum_vec, dp);
    }

    float temp[8];
    _mm256_storeu_ps(temp, sum_vec);
    for (int j = 0; j < 8; ++j)
    {
        result += temp[j];
    }

    for (; i < size; ++i)
    {
        result += row[i] * column[i];
    }

    return result;
}

```

Figure 11: Image of the dp_dpps.cpp file using machine vector instructions

```

Time to multiply two 16x16 float matrices: 9.08e-05 seconds
Time to multiply two 32x32 float matrices: 0.0002639 seconds
Time to multiply two 64x64 float matrices: 0.0007915 seconds
Time to multiply two 128x128 float matrices: 0.0045332 seconds
Time to multiply two 256x256 float matrices: 0.0294053 seconds
Time to multiply two 512x512 float matrices: 0.300061 seconds
Time to multiply two 1024x1024 float matrices: 3.83468 seconds

```

Figure 12: Execution times of various 2^n integer matrix sizes from $n = 16$ to $n = 1024$.




Figure 13: Graph measuring execution time vs. vector size using machine vector instruction DPPS code

Task#7: Compare all plots in one figure. Compare also to the performance plots from the previous take home test.




Figure 14: Graph comparing execution time vs. matrix size of compiler generated, optimized, and DPPS optimized assembly code




Figure 15: Graph comparing execution time vs. matrix size from previous take home test




Figure 16: Graph comparing execution time vs. matrix size from the previous take home test and current take home test

Conclusion:

The primary goal of this take-home test was to employ vector instructions to enhance dot product computations in the dp.cpp file. Utilizing vectorization commands in g++, along with DPPS instructions, I compiled and executed the files and observed that the runtime improved with increasing optimization levels. The slowest runtime was without any optimization, followed by optimization with assembly, and finally optimization with DPPS.